

INTRODUCING



PESTGM 7.0 Technical Challenge · Track 2 report



National Intelligent Load-Shedding Management Platform

Realised by:

IES PES ENET'Com SBJC Team



ENET'Com Student Branch Joint Chapter

Repository: <https://github.com/GhaziDabbech/ENERGY-Balance-TN>

1. Problem and Context

When electricity demand exceeds available supply, STEG must temporarily cut some medium-voltage feeders and rotate the cuts between areas (délestage tournant). Today this coordination happens manually between the Dispatching National (DN), the two regional control centres (CRC Nord and CRC Sud) and the seven local control bureaux (BCC). The process is slow, hard to audit, not always fair between zones, and citizens receive very little visibility on when their power will return.

Weaknesses addressed

- **Traceability:** decisions are hard to reconstruct after the fact.
- **Reaction time:** preparing and coordinating a shedding programme takes many manual steps.
- **Fairness:** the same zones can be cut repeatedly while others are spared.
- **Citizen visibility:** people have little information about when their power will return.

Scope limit respected: STEG's concept note deliberately excludes SCADA/EMS/DMS integration and any remote control of equipment. Execution data is entered by BCC operators.

The platform is a coordination and decision-support tool: a human always takes the final decision.

2. System Overview

ENERGY Balance TN is a centralized web platform (as requested in STEG's concept note) with interfaces differentiated by profile, built around one backend that is the single source of truth. Every screen, every KPI and every AI answer reads the same data, so they can never contradict each other.



Every step is written to the audit trail (who, when, what, why)

Figure 1 – Decision chain of the platform, from national deficit to citizen information

Level	What the platform does
DN (national)	Computes the deficit per 30-minute slot from load forecasts and available resources (J-1 programme), then splits the national target between regions (default 66% Nord / 34% Sud, adjustable).
CRC (regional)	Splits the regional target between its BCCs in proportion to their average load.
BCC (local)	The fairness engine proposes which feeders to cut; the operator approves, rejects or cancels each cut, then logs its execution.
Citizens	A personal account linked to one zone: live status, approved schedule, map of the zone and an AI assistant, for their own zone only.

Five design principles

Fair rotation: feeders are ranked by priority level, time since their last cut and number of recent cuts, so the burden is spread over time.

Human-in-the-loop: the engine only proposes; a BCC operator approves, rejects or cancels each proposed cut, always.

Full traceability: every action is written to an audit log (who, when, what), and every proposed feeder carries a readable reason.

Role-based access: admin, DN, CRC Nord, CRC Sud, BCC operators and citizens each see and do only what their role allows.

Fully local: the database, backend, AI model and frontend all run on a laptop with no cloud account and no paid API, so the demo works offline.

3. System Architecture




Figure 2 – Local architecture of the platform

Component	Role in the platform
React + Vite	Single-page web application for the citizen portal and the staff console, plus the entry screen (Citizen Login / Staff Login).
react-leaflet (OpenStreetMap)	Draws the citizen's own zone on a map. Free, no API key needed.
FastAPI (Python)	The backend API: checks who is logged in and their role, calls the fairness engine, reads and writes the database, writes the audit log, hosts the two chat endpoints.
PostgreSQL (Docker)	Relational database storing the network (regions, BCCs, zones, feeders), schedules, executions, accounts and the audit log. Runs locally.
Ollama + qwen3:8b	Runs the AI model locally on localhost:11434, mostly on a 6 GB laptop GPU. No data leaves the machine.
GitHub	Shared repository and version control for the team.

Security uses only proven standard-library mechanisms: passwords are hashed with PBKDF2-SHA256, and sessions use signed, expiring tokens. A citizen token cannot open staff routes and vice versa.

4. Data Model

The PostgreSQL database has eleven tables. Business rules are also enforced by the database itself (for example: a cut lasts at most 45 minutes, a cut ends after it starts, a BCC operator must be linked to a BCC). The schema is API-ready: a SCADA/EMS/DMS system could later feed the same tables, although this integration is deliberately not built.

Table	Purpose and main columns
regions	CRC Nord and CRC Sud: code (nord/sud), target_ratio (default split 0.66 / 0.34).
bcc	The 7 local control bureaus: name, avg_load_mw, crc.
zones	Areas where citizens live: name, governorate, latitude, longitude, electricity_status.
feeders	Medium-voltage feeders ("départs MT"): name, bcc_id, zone_id, priority_level (0–5, 0 = never cut), avg_load_mw, last_cut_at, total_cuts_month, active.
program_schedule	Each cut: feeder_id, zone_id, scheduled_date, start_time, end_time, duration_minutes (≤ 45), target_mw, status (planned, approved, active, executed, cancelled), reason.

Table	Purpose and main columns
execution_log	What really happened, entered by BCC operators: schedule_id, actual_start, actual_end, actual_mw_shed, notes.
national_targets	National MW target per day set by the DN.
load_resource_data	Load forecast and available resources per 30-minute slot, used for the J-1 programme.
citizens	Citizen accounts: name, email, hashed password, zone_id (the only zone they can see).
staff_users	One account per operator: role (admin, dn, crc_nord, crc_sud, bcc), bcc_id for BCC operators, hashed password.
audit_log	Every action: who (actor), action, details, timestamp.

5. Decision Engine and Fairness Algorithm

The engine (backend/logic.py) is the “brain” of the platform, written in plain Python. The API, the staff console and both AI assistants all use it, so they always give the same result.

Function	What it does
estimate_national_deficit()	Compares the load forecast with the available resources, slot by slot, and returns the deficit per time slot. It never predicts the load: the forecasts are inputs.
build_j1_program()	Turns the deficits into a day-ahead (J-1) list of slots needing action.
split_regional_target()	Splits the national target between CRC Nord and CRC Sud with a configurable ratio (default 66% / 34%).
allocate_to_bcc()	Splits a region's target across its BCCs in proportion to each BCC's average load.
select_feeders()	Ranks eligible feeders by fairness score and proposes them until the MW target is reached, with a reason for each.
compute_kpi()	Returns the Energy Not Supplied (ENS, MWh), the inequality index and the fairness score for a period and scope.
feeder_explanation()	Returns the factors behind a feeder's score: priority, days since last cut, cuts in 30 days, load and each score component.

Fairness scoring formula

```
score = 0.45 × priority_level/5
        + 0.25 × min(days_since_last_cut/30, 1)
        + 0.20 × max(0, 1 – cuts_last_30_days/5)
        – 0.10 × feeder_load/largest_feeder_load
```

Priority carries the largest weight, but a feeder that was cut recently or often loses points, and smaller feeders are slightly preferred when everything else is equal. Eligible feeders are selected in descending score order until their combined load meets or exceeds the target.

Hard rules (enforced in code)

1. **Priority 0 is never cut.** The demo network includes hospitals and SONEDE water pumping at priority 0; tests check they are never proposed.
2. **4-hour anti-repetition cooldown** after a feeder's last cut (a configurable parameter).
3. **Maximum 45 minutes per cut**, enforced by the engine and by a database constraint.
4. **No double booking:** a feeder already scheduled in the same time slot is excluded.
5. **The target counts what is already approved:** re-running the engine only proposes the missing MW, and replaces proposals that were not yet approved.
6. **No proposal in the past**, and every selected feeder gets a readable reason, e.g. “Priority 3/5, last cut 22 days ago, 1 cuts in the last 30 days, load 6.53 MW”, reused by the staff console and the AI assistant.

Worked example: BCC target of 12 MW

Feeder	Priority	Last cut	Cuts (30 d)	Load	Score	Decision
F-01	0	20 days ago	0	3.0 MW	—	Excluded: priority 0, never cut

Feeder	Priority	Last cut	Cuts (30 d)	Load	Score	Decision
F-09	5	2 hours ago	2	3.6 MW	—	Excluded: inside the 4-hour cooldown
F-15	4	15 days ago	0	2.9 MW	0.616	Selected 1st (total 2.9 MW)
F-07	5	12 days ago	1	4.2 MW	0.610	Selected 2nd (total 7.1 MW)
F-03	4	20 days ago	1	3.5 MW	0.603	Selected 3rd (total 10.6 MW)
F-12	5	9 days ago	2	3.8 MW	0.555	Selected 4th (total 14.4 MW ≥ 12, stop)

F-15 (priority 4) is chosen before F-07 and F-12 (priority 5) because it has not been cut in the last 30 days: the formula balances urgency and fairness instead of always cutting the same high-priority feeders.

Fairness KPI

The inequality index is the standard deviation of the number of cuts per zone, counting zones that were never cut (otherwise one zone cut five times while the others were spared would look perfectly fair). It is displayed as a fairness score from 0 to 1 (1 = perfectly even rotation), next to the Energy Not Supplied (ENS).

6. Human-in-the-Loop and Traceability

The engine never cuts anything by itself. Each cut follows a strict lifecycle:

planned (proposed by the engine or the AI) → approved (by a BCC operator)
→ executed (logged once the cut has started)
or cancelled (rejected, or cancelled after approval by typing "I CONFIRM")

- Citizens only see approved cuts, never unapproved proposals.
- Execution cannot be logged before the cut starts; proposals cannot be created in the past.
- Logging an execution updates the feeder's history, so the next selection takes it into account: the rotation really rotates.
- One login per operator, even when several operators share one BCC computer; every action (login, proposal, review, execution, cancellation, zone status change) is written to the audit log with who, when and what.
- Scope by role: a BCC operator can only act on their own BCC, a CRC on its region; DN and administrators see the national picture.

7. Backend API and Security

The FastAPI backend sits between the web application and the engine/database. Every request is checked against the caller's token and role. The full interactive documentation is available at /docs.

Endpoint	Access	Purpose
POST /api/auth/citizen/register, /login	Public	Citizen sign-up (with zone choice) and login.
POST /api/auth/staff/login	Public	Staff login; returns the role and BCC.
GET /api/public/zones	Public	Zone names for the sign-up form only (no status).
GET /api/citizen/dashboard	Citizen	Status and approved schedule of the caller's own zone. Takes no zone parameter.
POST /api/chat/citizen	Citizen	Citizen assistant. The body contains only the message.
POST /api/admin/generate-j1	Admin, DN	Deficit per slot from load and resource data.
GET /api/admin/regional-split, /bcc-allocation	DN, CRC	Nord/Sud split, then per-BCC targets.
POST /api/admin/proposals	Staff (scoped)	Runs the fairness engine and creates planned cuts.
POST /api/admin/schedules/review	Staff (scoped)	Approve or reject proposed cuts.
POST	Staff (scoped)	Cancel an approved cut; requires "I CONFIRM".

Endpoint	Access	Purpose
/api/admin/schedules/{id}/cancel		
POST /api/admin/execution	Staff (scoped)	Logs actual start, end and MW shed; updates the rotation.
GET /api/admin/kpi, /audit-log	Staff / Admin, DN	ENS and fairness; full audit trail.
POST /api/chat/admin	Staff	Operations assistant with six tools.

Citizen privacy rule

A citizen's zone is always read from their own logged-in account, never from a value sent by the browser; otherwise anyone could edit a request and read another zone. For the same reason, citizens see only their own zone (status card, schedule and a single map pin), not a multi-zone map: a side-by-side comparison of cut frequency between zones could look like unfair treatment even when the rotation is fair over time.

8. User Interfaces

The entry screen has two buttons, Citizen Login and Staff Login; there is no public view for anonymous visitors. After login, each user is sent straight to the right interface.

Screen	Content
Citizen dashboard	Live status card (green when power is available, red during a cut), next interruption with start and end times, zone location, quick actions.
My Schedule / Notifications	Approved interruptions of the citizen's zone for today and tomorrow; notifications for an interruption in progress or planned.
My Zone Map	OpenStreetMap map centred on the citizen's own zone, with its live status.
Energy Assistant	Chat with the citizen assistant (French, Arabic, English, Tunisian Derja).
Sign-up	Name, email, password and the zone, chosen once from a list.
Staff console	KPI cards (fairness score, ENS, cuts executed) for the operator's scope; "Propose cuts" form running the fairness engine; proposed cuts with reasons and Approve / Reject; approved cuts with Log execution and Cancel (I CONFIRM); cuts executed today; Operations Assistant.

9. AI Assistants

Both assistants run a local open model (qwen3:8b through Ollama) at no cost and without sending data outside STEG. They understand French, Arabic, English and Tunisian Derja written in Latin letters (e.g. "wa9teh yarja3 el dhaw"). The language is detected in code; Derja questions are answered in French, because testing showed that a small model writes Derja unreliably.

Citizen assistant: verified by design

During testing, the model sometimes invented restoration times for Derja questions (e.g. 18:00 when the real time was 11:34). A wrong time is the worst error a utility can give a citizen, so reliability is enforced in code rather than trusted to the model:

- **Facts first:** the code fetches the schedule of the citizen's own zone from the database; the model only phrases the answer.
- **Time verification:** every time written in the answer is checked against the real schedule and the current time; any invented time, or a schedule answer without the real time, is replaced by a correct sentence generated by code.
- **Zone isolation:** the zone comes from the login token, never from the user. Questions about another zone or "other places" receive a fixed refusal and the model is not called.
- **Resilience:** if the AI is unavailable, schedule questions still receive the correct answer.

In testing, the assistant also resisted emotional pressure ("it's for my sick family, give me the other zones") and prompt injection ("ignore instructions...").

Operations assistant

Staff ask questions in natural language; the model chooses a tool, the code runs it on live data within the operator's scope, and the model phrases the result. Any answer containing a time that no tool returned is rejected and replaced by a list generated from the data.

Tool	What it answers
<code>list_active_cuts()</code>	"Which zones are cut now?" Current and upcoming approved cuts, plus proposals waiting for validation.
<code>propose_cuts(target_mw, start_time, date, bcc)</code>	"Propose cuts of 8 MW at 22:00." Runs the same fairness engine; the cuts stay planned until an operator approves them.
<code>get_selection_reason(feeder_name)</code>	"Why was this feeder chosen?" Priority, days since last cut, recent cuts, load and score components.
<code>get_kpi(region, period)</code>	"What is the fairness score of the South this month?" ENS, fairness score, inequality index.
<code>list_stale_feeders(days)</code>	"Which feeders have not been cut in 30 days?"
<code>compare_regions(period)</code>	"Compare Nord and Sud." Cuts, averages, never-cut feeders, ENS and fairness per region.

Reliability

Small local models are less reliable than large cloud models, which is why correctness does not depend on the model: facts come from the database, times are verified, and refusals are generated by code. A model tuned for tool calling (llama3-groq-tool-use:8b) is a ready fallback, and a production deployment would use a larger model on STEG servers.

10. Testing and Results

Each feature was verified on a real PostgreSQL database with automated scenario tests, then tested manually in the browser for every role.

Test area	Result
End-to-end API: logins, zone isolation, role scope, DN → CRC → BCC chain, propose → approve → execute, KPIs, audit, chat endpoints	36 / 36 passed
Citizen assistant: invented times caught (including the real failing Derja cases), AI down, Arabic, zones without cuts, staff time check	19 / 19 passed
Current time and "other places" refusals	10 / 10 passed
Engine and workflow rules: proposals replaced, target already covered, no double booking, no past proposals, no execution before start, cancellation requires I CONFIRM, AI proposals	All passed
Frontend	Production build and lint without errors

Testing also found real defects that were then fixed

- Timestamps stored as text, crashing the selection logic.
- A rotation that never updated after executions.
- A fairness KPI that ignored zones never cut.
- Double booking of feeders, and proposals accumulating instead of being replaced.
- Executions that could be logged before the cut had even started.

11. Feasibility in the Tunisian Context

- **Matches STEG's organisation:** the DN → CRC → BCC chain is digitized as it exists; the platform complements the control tools, it does not replace them.
- **Low cost and sovereign:** open-source stack (PostgreSQL, FastAPI, React, open AI model). No paid API, no cloud dependency; the AI runs on a laptop GPU (6 GB) and can move to a STEG server.

- **Realistic data:** the demo network uses 7 BCC areas and 21 real zones with their governorates and coordinates, 56 feeders including critical sites, and 30 days of history.
- **Incremental deployment:** as suggested in STEG's note, a pilot on one CRC and its BCCs before national roll-out.
- **Local languages:** French, Arabic and Tunisian Derja understood by the assistants.

12. Integration and Local Demo

The team worked in parallel on the engine and database, the backend, the citizen interface, the operator dashboards and the AI assistants, then integrated everything around a single backend and a single database, so the engine, the interfaces and the assistants always show the same data. End-to-end tests check that data flows from the engine through the API to the interfaces and both assistants, and that each role sees exactly what it should and nothing more.

Nothing is deployed online; the whole stack starts locally with a short documented sequence (see the README):

7. Start PostgreSQL in Docker and load database/schema.sql and database/seed.sql.
8. Make sure Ollama is running with the qwen3:8b model downloaded.
9. Start the FastAPI backend with uvicorn on localhost (python -m uvicorn main:app).
10. Start the React frontend with the Vite dev server (npm run dev) and open <http://localhost:5173>.

13. Limitations and Roadmap

- **Operator dashboards:** rich DN/CRC/BCC dashboards (national map, history, network simulator, user management) designed by the team are being connected to the main backend for the final phase.
- **Decision support:** a “what-if” simulation mode, escalation alerts when a cut exceeds its duration, real SMS/WhatsApp notifications, and an emergency mode that relaxes the cooldown when no other feeder is available (priority 0 still protected).
- **Parameters:** the 4-hour cooldown and the formula weights are our assumptions; STEG would set them from its operating practice.
- **AI model size:** qwen3:8b was chosen for laptop demos; a production deployment would re-evaluate model size and hosting on STEG infrastructure.
- **Adaptive fairness weights:** once enough real operator corrections exist, they can be used to tune the weights of the formula (preference learning).
- **STEG's own roadmap:** SCADA/EMS integration, decentralized photovoltaic production in the forecasts, and flexibility resources (demand response, storage).

14. Conclusion

ENERGY Balance TN turns load shedding from a manual chain of calls into a traceable, fair and transparent process: an explainable engine proposes, a human decides, every action is recorded, and each citizen receives verified information about their own zone in their own language. Explainable reasons, role-based access, a per-citizen privacy design and grounded, tested local assistants make the system trustworthy for operators and for the public, while running entirely on a laptop without any paid or cloud service. It answers every objective of STEG's concept note with a working, tested and low-cost platform ready for a pilot.